





Token Efficiency



= 4 tokens

The AI Bill Surprise



AI accounting = magic trick gone wrong



BURNOUT TRACKER



TO DO LIST

- ☒ Overworked
- ☒ Stressed
- ☐ Engaged

IN PROGRESS

- Task 1
- Task 2
- Task 3
- Task 4
- Task 5

Project Setup

- devcontainer.json
- Copilot instructions
- **A**gent profile




Reduce input and context tokens

- Drop unused tools
- Scope sessions tightly
- Keep prompts precise and short



Choose the right execution path

- Use the smallest model that fits
- Prefer a direct completion
- Watch output, not just input





Token Efficiency

Run an instant model call, compare prompt cache reuse, then compact long working notes into a smaller reusable prompt.

MODEL
gpt-chat-latest

METER
5.5 ShortCo

REGION
westus3 / GL

LIVE PRICING
Instant Demo

An **instant model** is called by name — there's no deployment to create or manage, and it's available to every Foundry project automatically. It bills at standard pay-as-you-go rates, so it stays much cheaper than a data-zone or priority deployment. This call is priced live against Azure Retail Prices meters, then scaled to 1,000 and 1M runs.

1 call

Prompt

In one sentence, why are instant models useful for prototyping?

REUSABLE PREFIX
Prompt Cache Demo

Watch the model cache warm up in real time. A warm-up call primes a long stable prefix, then the repeated call reuses it — see the cache fill, the cold→warm comparison, and exactly what was loaded.

2 calls

Warm the cache

CONVERSATION HYGIENE
Compaction Demo

Long working notes pile up tokens on every turn. One call compacts them into a short, durable summary that keeps only the next-turn essentials — goal, key facts, and commands — so future prompts stay cheap. See the before→after token drop and reduction rate.


1 call

Working notes

Goal: ship a small Java dashboard change that explains instant model throughput limits.
Context gathered: Microsoft Foundry instant models are preview-scoped to West US 3 and can be called by model


Instant models: no deployment

- Call gpt-chat-latest by name, instantly
- No reserved throughput, no endpoint config
- Up to 2.5x cheaper than deployed tiers




Prompt caching: reuse a stable prefix

- Send the same long context twice
- Cache hit read straight from the model
- 96% hit slashes the repeat cost



Compaction: shrink long sessions

- Summarize history into durable context
- Spend tokens now, save every later turn
- 419 tokens become a 265-token summary



Demo

<https://aka.ms/costs>

